

# Migrating Ingress Controllers: The Six-Month Journey

**Martin Beránek - ShipMonk's Traefik  
evangelist**

**shipmonk**





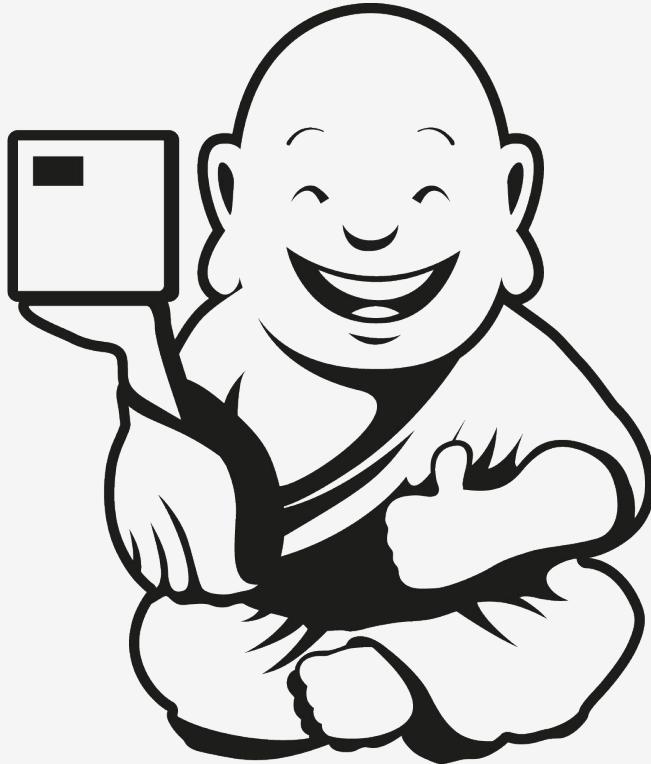
# Introduction



# The history of ingress setup

## Introduction

- ~2019 introduction of K8s instead of DO droplets
- Routing done through AWS ALB
  - Used AWS Cognito for private services
- Later changed to NLB with Nginx
- It worked!
  - Minor issues with latencies
  - Configs in configmap
  - Access logs - disabled after a while





# Retrospective

## Introduction

- NGINX does not offer [good] metrics
  - NGINX plus does
- High utilization under high connection count
  - Caused outage during the peak
  - No action item, nothing in logs
- Upstream setup causing unknown issues
  - Which was there cause of http 503

## Service Upstream

By default the Ingress-Nginx Controller uses a list of all endpoints (Pod IP/port) in the NGINX upstream configuration.

The `nginx.ingress.kubernetes.io/service-upstream` annotation disables that behavior and instead uses a single upstream in NGINX, the service's Cluster IP and port.

This can be desirable for things like zero-downtime deployments . See issue [#257](#).

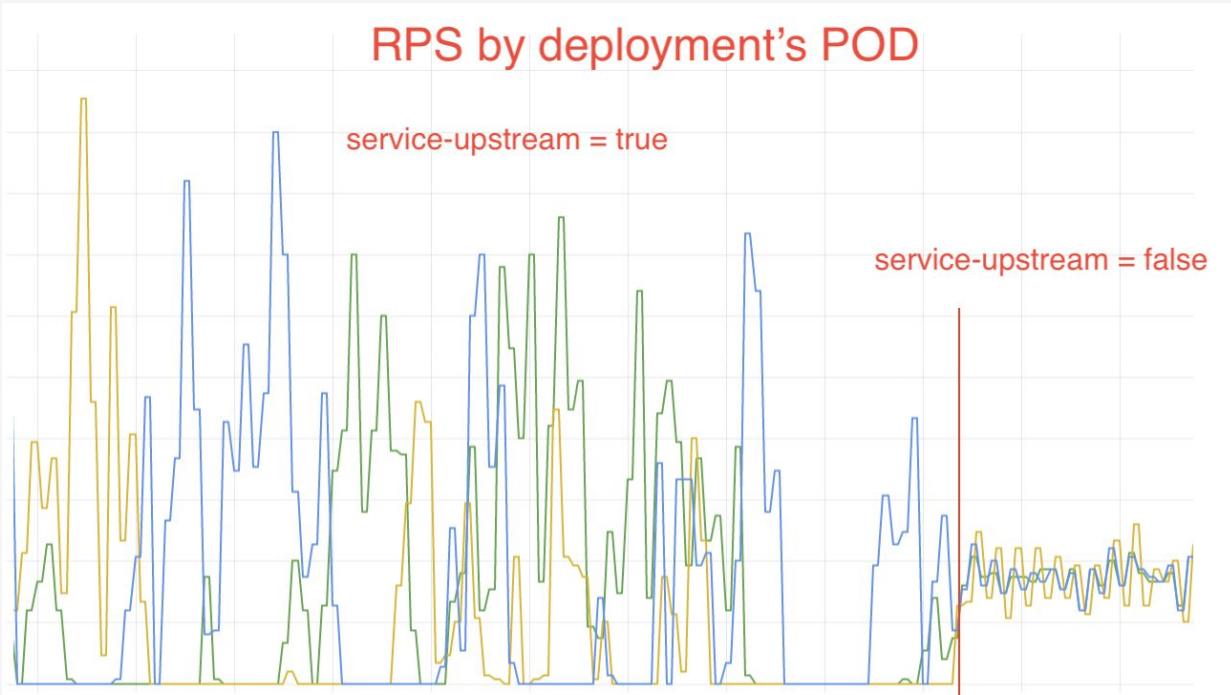
### Known Issues

If the `service-upstream` annotation is specified the following things should be taken into consideration:

- Sticky Sessions will not work as only round-robin load balancing is supported.
- The `proxy_next_upstream` directive will not have any effect meaning on error the request will not be dispatched to another upstream.

## RPS by deployment's POD

shipmonk



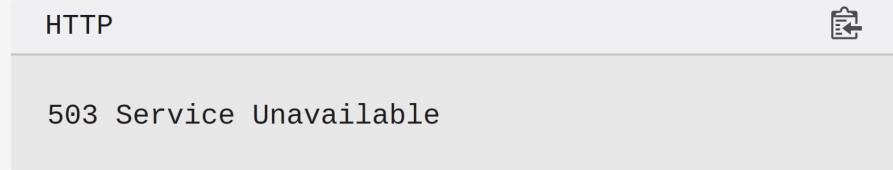
Issue #257, @narqo [comment](#)



# Upstream setup causing unknown issues

## Introduction

- To avoid HTTP 503 we switched service-upstream=true
- After 2-3 months, we noticed that backends were not utilized evenly
- Overall making HPA unusable and creating more HTTP 503

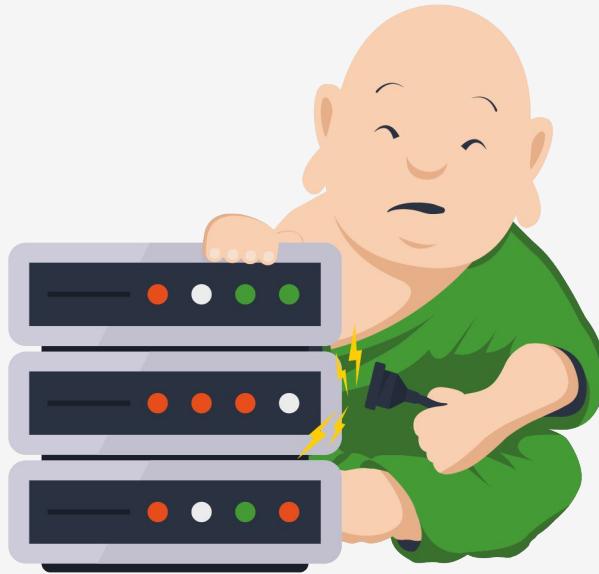




# Requirements

## Introduction

- Use a similar setup, no new CRDs
- Have the same setup on production and dev
- Use something which the team already partially knows
- It has to have [at least] metrics in Prometheus
- Be stable





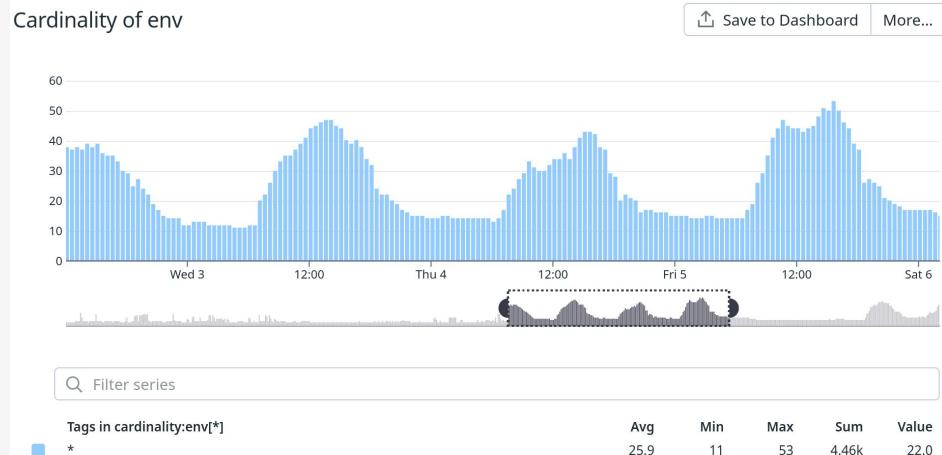
# Devs

## Introduction

- Anyone can deploy from their own MR
- Environments live for four hours
  - Can be resurrected by redeploy
- Each environment contain about 8 ingresses
  - Overall 32 rules

The overall cost of AWS ALB [~700 USD](#) due to rule count

For AWS NLB only ~50USD



# Plan:

⌚ The whole load balancing setup

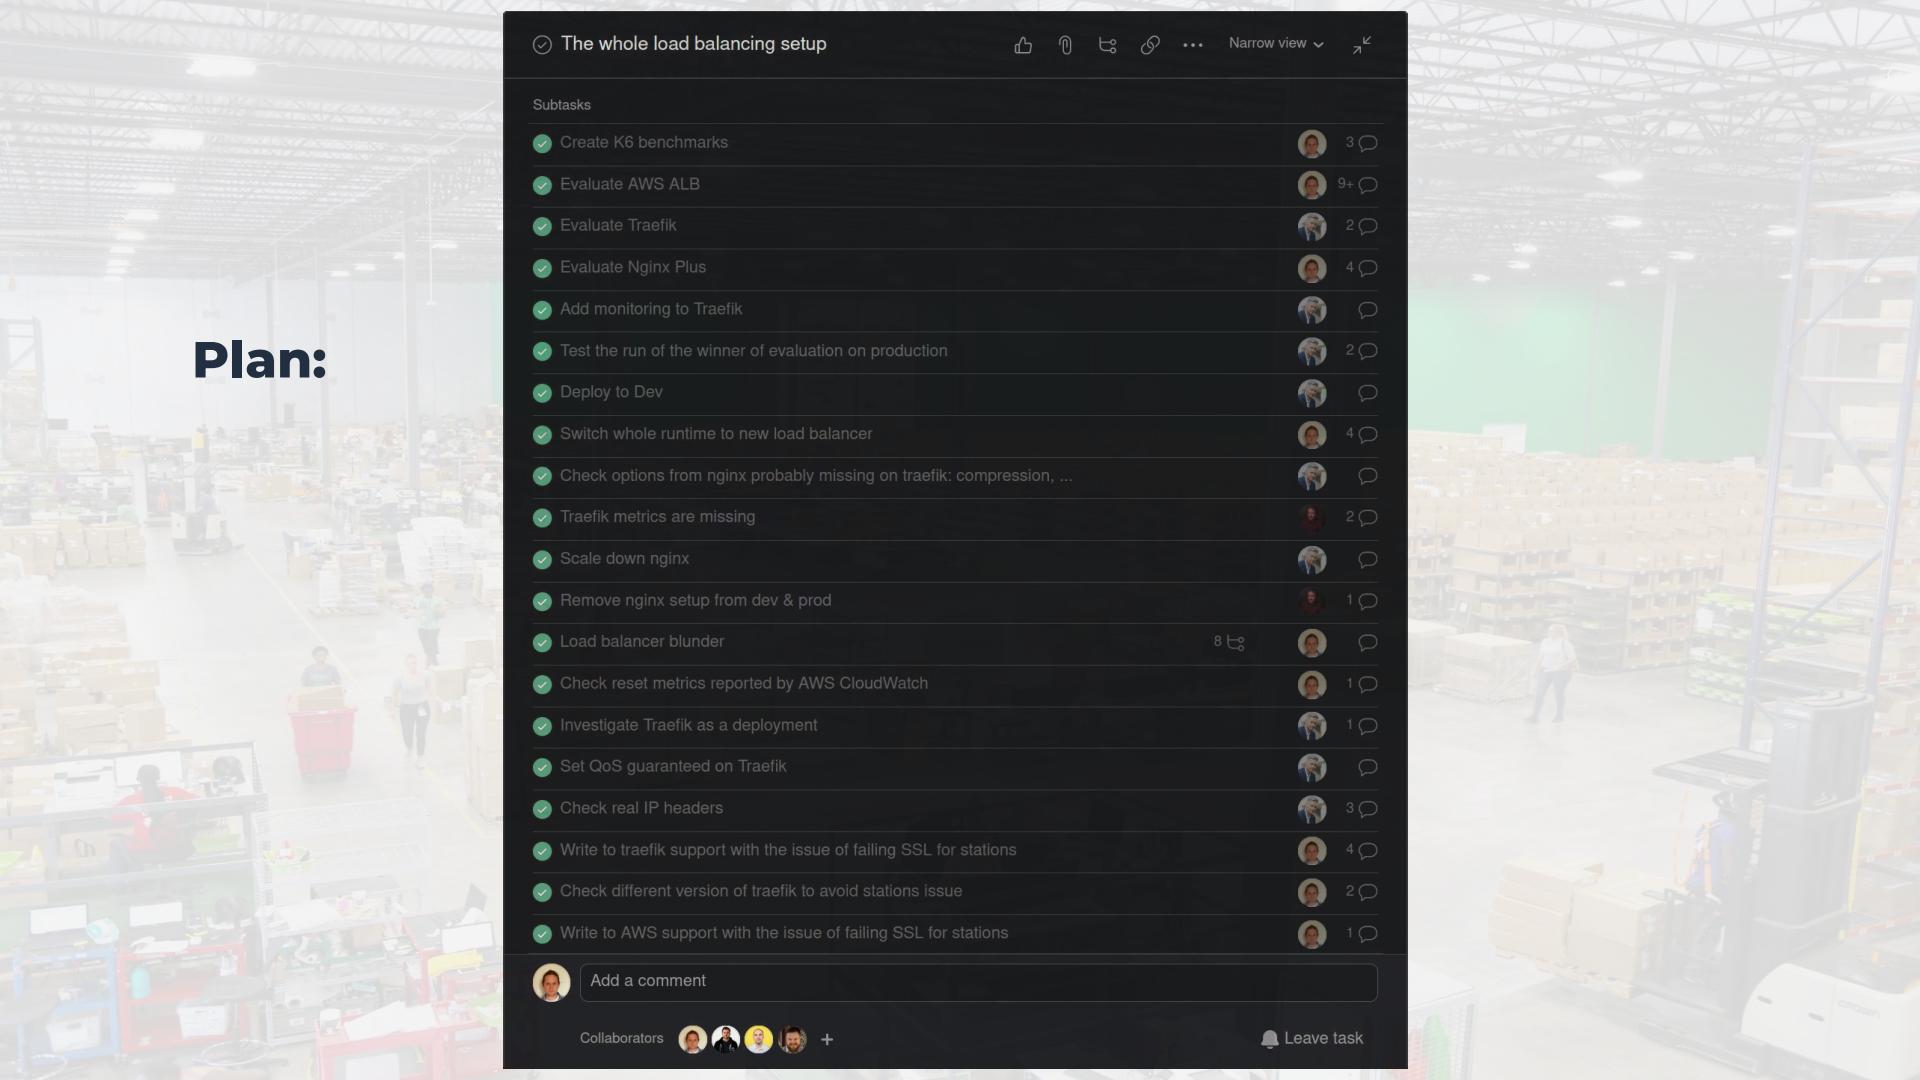
Subtasks

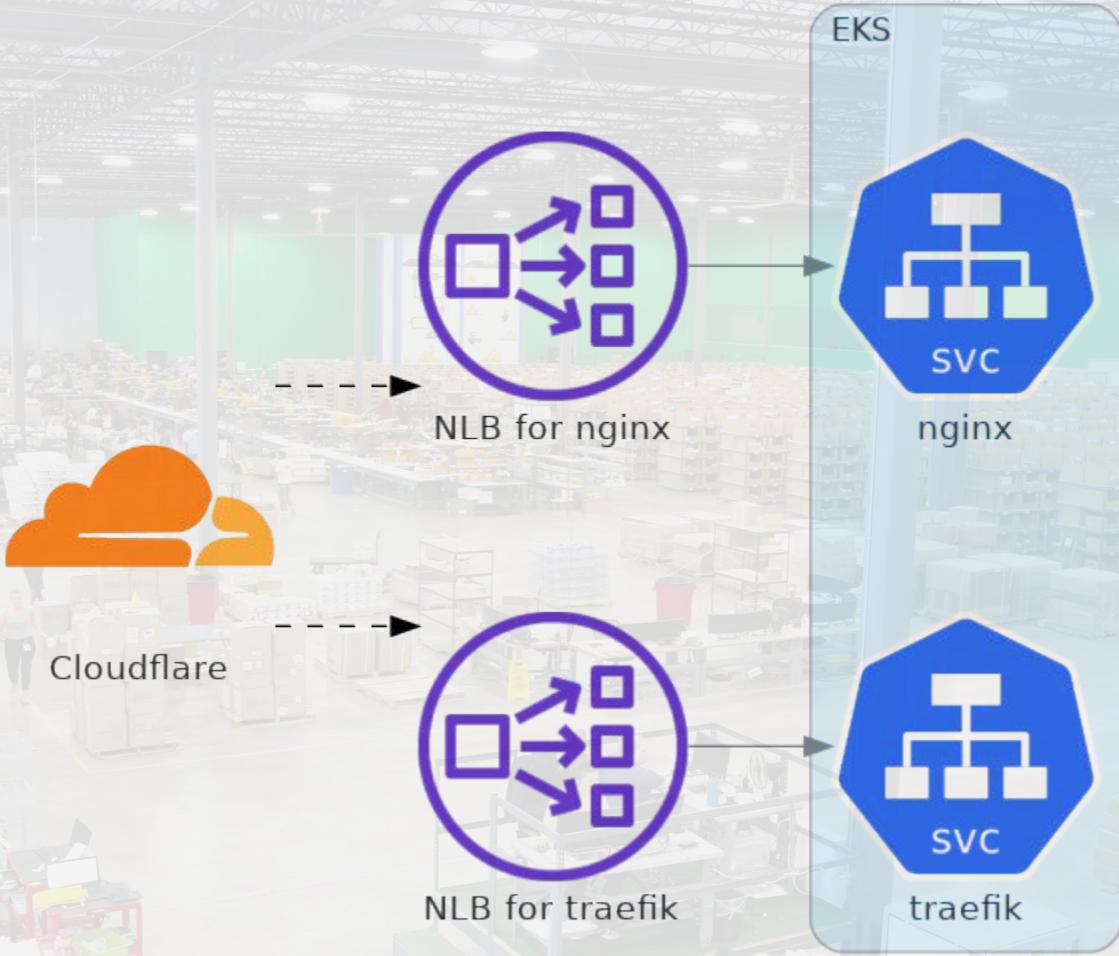
- ✓ Create K6 benchmarks
- ✓ Evaluate AWS ALB
- ✓ Evaluate Traefik
- ✓ Evaluate Nginx Plus
- ✓ Add monitoring to Traefik
- ✓ Test the run of the winner of evaluation on production
- ✓ Deploy to Dev
- ✓ Switch whole runtime to new load balancer
- ✓ Check options from nginx probably missing on traefik: compression, ...
- ✓ Traefik metrics are missing
- ✓ Scale down nginx
- ✓ Remove nginx setup from dev & prod
- ✓ Load balancer blunder
- ✓ Check reset metrics reported by AWS CloudWatch
- ✓ Investigate Traefik as a deployment
- ✓ Set QoS guaranteed on Traefik
- ✓ Check real IP headers
- ✓ Write to traefik support with the issue of failing SSL for stations
- ✓ Check different version of traefik to avoid stations issue
- ✓ Write to AWS support with the issue of failing SSL for stations

Add a comment

Collaborators

Leave task







# What took us so long?

- TLS suites issue 0-1 days
- EPROTO error in Nodejs 80 days
- Non TLS traffic on TLS endpoint 120 days
- Usual suspect: Traefik  $\infty$  days

# Partners are old

## TLS suites issues

- Most of the times, you are cooperating with old, very old software
- ICM\_HTTP\_SSL\_ERROR connecting to the app.shipmonk.com
- Failed connection to Locus robots
- Failed to connect to SAP



# Solution? Downgrade

## TLS suites issues



### Cipher Suites

#### # TLS 1.3 (suites in server-preferred order)

[TLS\\_AES\\_128\\_GCM\\_SHA256 \(0x1301\)](#) ECDH x25519 (eq. 3072 bits RSA) FS

[TLS\\_AES\\_256\\_GCM\\_SHA384 \(0x1302\)](#) ECDH x25519 (eq. 3072 bits RSA) FS

[TLS\\_CHACHA20\\_POLY1305\\_SHA256 \(0x1303\)](#) ECDH x25519 (eq. 3072 bits RSA) FS

#### # TLS 1.2 (suites in server-preferred order)

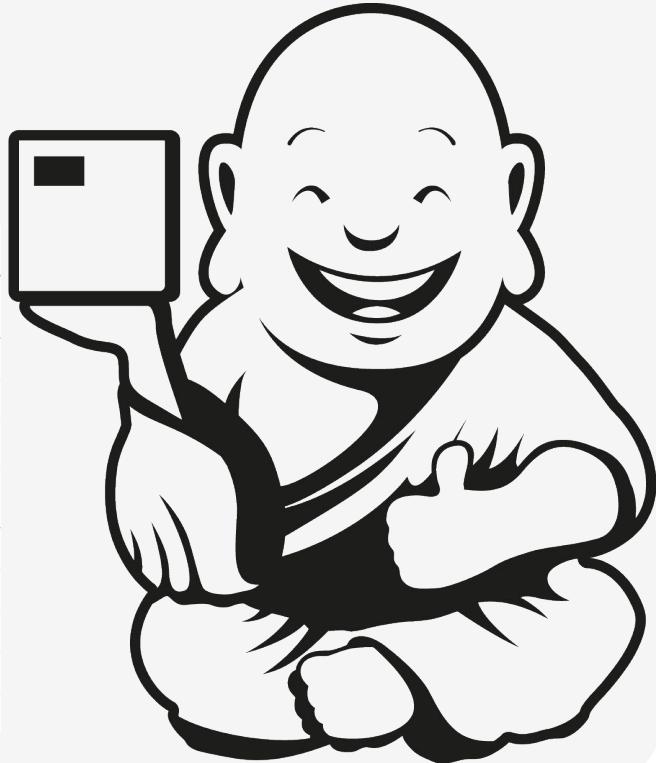
[TLS\\_ECDHE\\_RSA\\_WITH\\_CHACHA20\\_POLY1305\\_SHA256 \(0xccaa8\)](#) ECDH secp521r1 (eq. 15360 bits RSA) FS

[TLS\\_ECDHE\\_RSA\\_WITH\\_AES\\_128\\_GCM\\_SHA256 \(0xc02f\)](#) ECDH secp521r1 (eq. 15360 bits RSA) FS

[TLS\\_ECDHE\\_RSA\\_WITH\\_AES\\_256\\_GCM\\_SHA384 \(0xc030\)](#) ECDH secp521r1 (eq. 15360 bits RSA) FS

[TLS\\_ECDHE\\_RSA\\_WITH\\_AES\\_128\\_CBC\\_SHA256 \(0xc027\)](#) ECDH secp521r1 (eq. 15360 bits RSA) FS **WEAK**

(P) This server prefers ChaCha20 suites with clients that don't have AES-NI (e.g., Android devices)





# Fixed in 1 day

**TLS suites issues**

# Proxy protocol v2

## EPROTO error in Nodejs

- Protocol for transporting [connection information](#)
- Should it be on or off?

# Let's keep it on!

### Target unhealthy state management

Terminate connections on unhealthy targets

On

### Target deregistration management

Terminate connections on deregistration

Off

Deregistration delay (draining interval)

120 seconds

### Traffic configuration

Proxy protocol v2

On

Preserve client IP addresses

On

### Target selection configuration

Stickiness

Off

Cross-zone load balancing

Inherit settings from Network Load Balancer

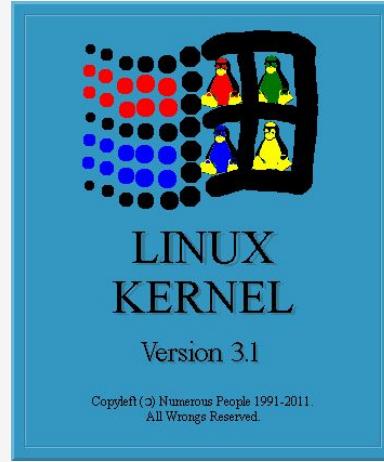


```
[24. 5. 2023 8:30:05] Request failed: FetchError: request to https://shipmonk-dev-eu-west-1-tfk-pub-6526a9003aa92c35.elb.eu-west-1.amazonaws.com/api/v1/wms/orders/images/refresh-from-remote failed, reason: write EPROTO 139976298616768:error:1408F10B:SSL routines:ssl3_get_record:wrong version number:../deps/openssl/openssl/ssl/record/ssl3_record.c:332:  
request to https://shipmonk-dev-eu-west-1-tfk-pub-6526a9003aa92c35.elb.eu-west-1.amazonaws.com/api/v1/wms/orders/images/refresh-from-remote failed, reason: write EPROTO 139976298616768:error:1408F10B:SSL routines:ssl3_get_record:wrong version number:../deps/openssl/openssl/ssl/record/ssl3_record.c:332:  
EPROTO
```

# Warehouse stations sentry issue

## EPROTO error in Nodejs

- Frontend developers reported SSL issues from nodejs app
- The app is very old
  - Old SSL/TLS suite?
  - Old OS?
  - ...
- This can't be us, can be?

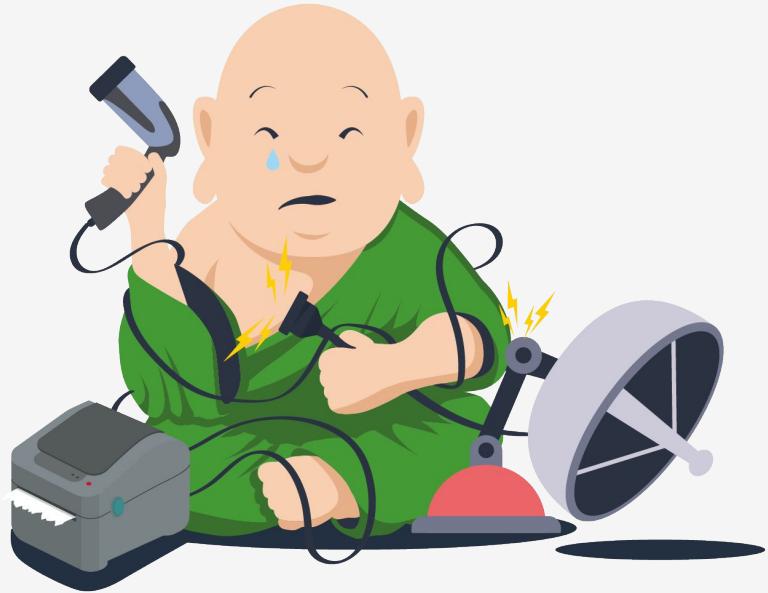


[Linux Kernel 3.1 proposed logo](#)

# Error replication

## EPROTO error in Nodejs

- Using the same node version, the same packages replicates the error
- Using the same code in different languages does not cause the issue



```

const fetch = require('node-fetch');
const https = require('https');

const url      = "https://shipmonk-dev-eu-west-1-tfk-
pub-65269e03aa92c35.elb.eu-west-1.amazonaws.com/api/v1/wms/orders/images
/refresh-from-remote"
const hostHeader = "dev30783.shipmonk.dev"

let running = true;

process.on('SIGTERM', () => {
  running = false;
});

async function run() {
  while (running) {
    try {
      const res = await fetch(url, {
        method: 'POST',
        headers: {
          'Content-Type': 'application/json',
          Host: hostHeader,
        },
        port: 443,
        agent: new https.Agent({ rejectUnauthorized: false }),
        body: JSON.stringify({
          fulfillmentId: 51292698,
        })
      });

      const date = new Date().toLocaleString();
      const status = res.status;
      const body = await res.text();
      if (!([202, 400].includes(status)))
        console.log(`[${date}] Response status: ${status}\nBody:
${body}`);
    } catch (err) {
      console.error(`[${new Date().toLocaleString()}] Request failed:
${err}\n${err.message}\n${err.code}`);
    }
  }
}

for (let i = 0; i < 10; i++)
  run();

```

# Nodejs → Golang



~60 lines

```

package main

import (
  "net"
  "crypto/tls"
  "encoding/json"
  "fmt"
  "io"
  "net/http"
  "os"
  "os/signal"
  "sync"
  "time"
)

const {
  URL      = "https://shipmonk-th-dev-pub.shipmonk.dev/api/v1/wms/orders/images/refresh-from-remote"
  hostHeader = "dev30783.shipmonk.dev"
  interval = // time between requests in seconds
}

var {
  running = true
  wg = sync.WaitGroup{
  }
}

type request struct {
  fulfillmentID string `json:"fulfillmentId"`
  Apisecret     string `json:"apisecret"`
}

func main() {
  signal.Notify(os.Signal(SIGPOLL), os.Interrupt)
  go func() {
    signal.Notify(os.Signal(SIGTERM), os.Interrupt)
    go func() {
      signal.Notify(os.Signal(SIGINT), os.Interrupt)
      running = false
    }()
  }()
  for i := 0; i < 100; i++ {
    go run()
  }
  wg.Wait()
  fmt.Println("Exiting...")
}

func run() {
  defer wg.Done()
  Apisecret := os.Getenv("APP_SECRET")
  transport := &http.Transport{
    TLSClientConfig: &tls.Config{InsecureSkipVerify: true},
  }
  client := &http.Client{Transport: transport}
  for {
    req, err := http.NewRequest("POST", URL, bytes.NewReader([]byte(`{"fulfillmentId": "51292698", "apisecret": "` + Apisecret + `"}`)))
    if err != nil {
      if err.In(ErrMalformedRequest, ErrRequestFailed) {
        continue
      }
      req, err := HTTP.NewRequest("POST", URL, bytes.NewReader(reqBody))
      if err != nil {
        fmt.Printf("%s\nRequest failed: %v", time.Now().Local().String(), err.Error())
        continue
      }
      req.Header.Set("Content-Type", "application/json")
      req.Host = hostHeader
      req.Method = "POST"
      if err := req.DoReq(); err != nil {
        if err.In(ErrMalformedRequest, ErrRequestFailed) {
          continue
        }
        defer resp.Body.Close()
        data := time.Now().Local().String()
        if resp.StatusCode != 202 {
          body := resp.ReadAll(resp.Body)
          if err := json.Unmarshal(body); err != nil {
            fmt.Printf("%s\nError reading response body: %v\n", time.Now().Local().String(), err.Error())
            continue
          }
          fmt.Printf("%s\nResponse status: %d\nBody: %s\n", data, resp.StatusCode, string(body))
        }
      }
    }
  }
}

```

~100 lines

**“... that's node problem,  
let's ignore it”**



# Fixed in 80 days

**EPROTO error in Nodejs**

**“A few days later...”**



**Robert Josiek (CZ, WMS Product Manager)** 6 months ago  
<https://metabase.shipmonk.cloud/> down again



55 replies



**Richard Felkl (CZ DevOps engineer)** 6 months ago  
not for me



**Tomáš Hejátko (CZ DevOps Engineer)** 6 months ago  
works for me too



**Robert Josiek (CZ, WMS Product Manager)** 6 months ago  
image.png ▾



This site can't provide a secure connection

metabase.shipmonk.cloud sent an invalid response.

ERR\_SSL\_PROTOCOL\_ERROR

Reload

# Traefik returned non TLS content on TLS endpoint

## Non TLS traffic on TLS endpoint

- What was that content about?
  - What the Wireshark session got me:



```
HTTP/1.1 400 Bad Request
Content-Type: text/plain; charset=utf-8
Connection: close
400 Bad Request
```

# What we knew?

## Non TLS traffic on TLS endpoint

- Everytime we switched 100% of traffic to Traefik, issue started after 16pm Prague time, 9am EST
- But we benchmarked everything with k6, is the load an issue?
- This never happened with 10% or 50% of traffic

**Let's bomb it till it starts to talk!**



- It's not easy to force app **not** reuse connections
- Even if they are not reused, they can be cleanup anytime
- Let's create the largest amount of concurrent connections we can
- Just never ...Body.Close()

```
● ● ●  
...  
  
func main() {  
    ...  
  
    var wg sync.WaitGroup  
  
    for i := 0; i < numConns; i++ {  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            for true {  
                req, err := http.NewRequestWithContext(traceCtx, http.MethodGet, url, nil)  
                if err != nil {  
                    log.Fatal(err)  
                }  
                _, err = httpClient.Do(req)  
                if err != nil {  
                    log.Fatal(err)  
                }  
                // _.Body.Close()  
            }  
        }()  
    }  
    wg.Wait()  
}
```

# Healthy



```
while true; do netstat --all --program  
2>/dev/null | grep `pgrep main` 2>/dev/null | wc  
-l; sleep 1; done
```

```
0  
0  
0  
555  
2994  
3279  
4540  
4540  
4540  
5127  
5518  
5518  
6110  
6405  
^C%
```

shipmonk

# Not Healthy



```
while true; do netstat --all --program  
2>/dev/null | grep `pgrep main` 2>/dev/null | wc  
-l; sleep 1; done
```

```
0  
0  
0  
555  
2994  
0  
0  
0  
^C%
```

# Solution? Dunno

## Non TLS traffic on TLS endpoint

We tried:

- Opening ticket to AWS
- Opening ticket to Traefik partner
- Rewrite ingress to IngressRoute (Traefik CRD)
- Different builds AMD64 instead of ARM
- Moving versions up and down
- Scaling turned off & on, pre scaled to ~100 of Traefics
- Opened an issue on GitHub

# Solution? Disable Proxy v2

Non TLS traffic on TLS endpoint

## Target unhealthy state management

Terminate connections on unhealthy targets

On

## Target deregistration management

Terminate connections on deregistration

Off

Deregistration delay (draining interval)

120 seconds

## Traffic configuration

Proxy protocol v2

Off

Preserve client IP addresses

On

## Target selection configuration

Stickiness

Off

Cross-zone load balancing

Inherit settings from Network Load Balancer



# Fixed in 180 days

**Non TLS traffic on TLS endpoint**

# It got fixed recently

Non TLS traffic on TLS endpoint



juliens commented 2 weeks ago

Member

...

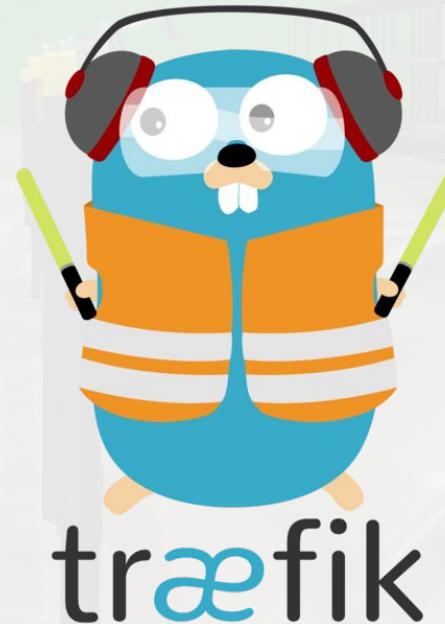
Upon conducting additional analysis, it appears that the issue is linked to a read header timeout. Upon further inspection, I discovered that this timeout was not configured in the PROXY protocol library we are utilizing. Consequently, I implemented a solution.

A big thank you to @jb-abbadie for providing a reproducible case; it was immensely helpful.

shipmonk



# Usual suspect: Traefik



traefik

## So can we confirm Traefik is causing it?

It works for me but not for [@Petr Morávek \(CZ SW Engineer\)](#), is **traefik** working correctly?

Thanks everyone for identifying the root cause

[@Martin Beránek \(CZ INF Eng. lead\)](#) this was another issue caused by the switch to **Traefik**. Can we implement a better process for catching similar errors?

You saw it in [#major-system-issues](#), correct?

cc [@Matěj Smišek \(CZ VP of Engineering\)](#)

This was another issue with **Traefik**, [@Martin Beránek \(CZ INF Eng. lead\)](#) we have to think of a way how to test these changes directly after they are done so we don't disrupt production and operations

I know but where can we have a discussions? Should we create a **traefik** possibly related issue temp channel?

shipmonk

[@Martin Beránek \(CZ INF Eng. lead\)](#) can we revert Traefik and check with them if it works again?

## Traefik was deployed yesterday, maybe it's causing issues?

Previously it was caused by the switch to **Traefik**, co you check if it's still the case?

Locus sync is unstable, not sure if **Traefik** would cause it, it's happening from last night

<https://shipmonk.slack.com/archives/C04GQ6FD01F/p1690355324640869>

Probably a short issue when removing **Traefik** <https://shipmonk.slack.com/archives/C03E989J37F/p1683014931176319>

 Richard Felkl (CZ DevOps engineer)

 [#channel](#) Today from 11AM to 12AM CET (5AM - 6AM EST / 2AM - 3AM PST) we're going to switch 50% of production load to Traefik load balancer again. There should be no outage. If any problem happens let us know to this thread.

Thread in [# devops\\_status](#) | May 2nd | [View message](#)

 3 replies

Can we turn off **traefik** balancing for the backend app and see what effect it has on the latency? <https://app.datadoghq.c...> Show more (edited)

Why would this happen? What does **Traefik** do differently? It shouldn't cause unreachable apps

# People tend to blame every issue on Traefik

**Usual suspect: Traefik**

- With 90% of accuracy
- But those 10% hurt a lot
- Even in cases there was only Nginx running
- Or after Traefik was running without issues for months

**The only cure is the overall stability**



# The flow

## Usual suspect: Traefik

- Evaluate multiple options (benchmark, check docu, ...)
- Announce the change as detailed as possible
- Deploy to reasonable amount of traffic: 10%, 50%, ...
- Evaluate, do reviews
- Announce the final deployment
- Deploy and evaluate again
- Cleanup of old setup





# Fixed in $\infty$ days

**Usual suspect: Traefik**



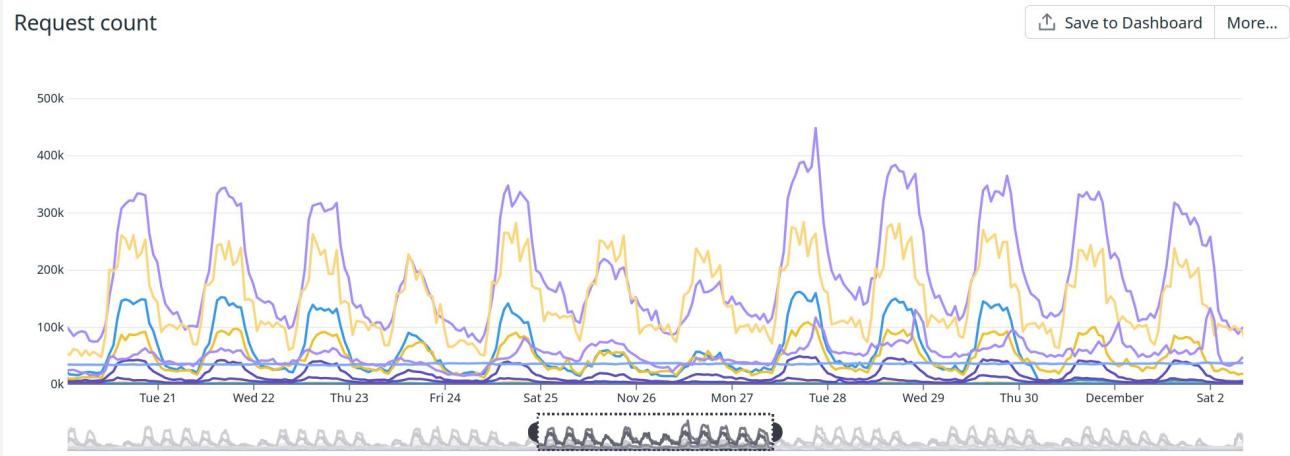
# Review

# 24.11.23 the Peak, Black Friday

## Review

shipmonk

- No utilization issues
- Only 4 replicas
- About 0.1 CPU load for each
- Around 1 GB of RAM
- Average day



We had ~10 minutes downtime last year cause of NGINX issues

# Sizing

## Review

- Original setup used on production
  - 8 pods
  - Each pod 1 CPU and 2 GiB
- Much larger scale on dev
  - Each pod 500 mCPU and 256 MiB
  - As DaemonSet
- Traefik has on prod:
  - 4 pods
  - Each pod 1.3 CPU and 2.4 GiB
- On dev:
  - 2 pods
  - Each pod 0.7 CPU and 1.2 GiB

# Rightsizing

## Review

shipmonk

### Summary

Type	Usage	Requests	Limits
CPU	8.02%	1.30 Cores	1.30 Cores
Memory	25.70%	2.34 GB	2.34 GB

### Recommendation

Type	Recommended Requests
CPU	0.17-0.21 Cores
Memory	1.00-1.20 GB



# Requirements

## Introduction

- Use a similar setup, no new CRDs
- Have the same setup on production and dev
- Use something which the team already partially knows
- It has to have [at least] metrics in Prometheus
- Be stable



# Thank You

Questions?

shipmonk



# We are hiring

Try our challenge in this docker image

**shipmonkdevops/shipmonk-test**

**shipmonk**

